

ECE 780 To 8 Project

Bug Algorithms using Neural Networks

Andreas Stöckel

Computational Neuroscience Research Group
Centre for Theoretical Neuroscience



July 24th, 2017

Motivation

- ▶ In class, we mostly talked about path planning using *global* information.
- ▶ Often, robotic and biological agents alike are restricted to *local*, *sensor-based* information.



Source: Wikimedia Commons

Motivation

- ▶ In class, we mostly talked about path planning using *global* information.
- ▶ Often, robotic and biological agents alike are restricted to *local*, *sensor-based* information.
- ▶ *Bug algorithms* (Lumelsky and Stepanov 1987) are a simplistic class of local planning algorithms/policies.



Source: Wikimedia Commons

Motivation

- ▶ In class, we mostly talked about path planning using *global* information.
 - ▶ Often, robotic and biological agents alike are restricted to *local*, *sensor-based* information.
 - ▶ *Bug algorithms* (Lumelsky and Stepanov 1987) are a simplistic class of local planning algorithms/policies.
- ~> Can these algorithms be implemented in a biological, neural substrate?



Source: Wikimedia Commons


PART I

Bug Algorithms

Bug Algorithm Assumptions

Environment

- ▶ Start location \vec{p}_{start} , goal location \vec{p}_{goal}
- ▶ Polygonal obstacles O_1, \dots, O_n




Bug Algorithm Assumptions

Environment

- ▶ Start location \vec{p}_{start} , goal location \vec{p}_{goal}
- ▶ Polygonal obstacles O_1, \dots, O_n

Robot

- ▶ Knows direction towards goal (absolute and relative)




Bug Algorithm Assumptions

Environment

- ▶ Start location \vec{p}_{start} , goal location \vec{p}_{goal}
- ▶ Polygonal obstacles O_1, \dots, O_n

Robot

- ▶ Knows direction towards goal (absolute and relative)




Bug Algorithm Assumptions

Environment

- ▶ Start location \vec{p}_{start} , goal location \vec{p}_{goal}
- ▶ Polygonal obstacles O_1, \dots, O_n

Robot

- ▶ Knows direction towards goal (absolute and relative)
- ▶ Knows the straight line distance to the goal



Bug Algorithm Assumptions

Environment

- ▶ Start location \vec{p}_{start} , goal location \vec{p}_{goal}
- ▶ Polygonal obstacles O_1, \dots, O_n

Robot

- ▶ Knows direction towards goal (absolute and relative)
- ▶ Knows the straight line distance to the goal
- ▶ Has no knowledge of obstacles (no map)




Bug Algorithm Assumptions

Environment

- ▶ Start location \vec{p}_{start} , goal location \vec{p}_{goal}
- ▶ Polygonal obstacles O_1, \dots, O_n

Robot

- ▶ Knows direction towards goal (absolute and relative)
- ▶ Knows the straight line distance to the goal
- ▶ Has no knowledge of obstacles (no map)
- ▶ Has a contact sensor




Bug Algorithm Assumptions

Environment

- ▶ Start location \vec{p}_{start} , goal location \vec{p}_{goal}
- ▶ Polygonal obstacles O_1, \dots, O_n

Robot

- ▶ Knows direction towards goal (absolute and relative)
- ▶ Knows the straight line distance to the goal
- ▶ Has no knowledge of obstacles (no map)
- ▶ Has a contact sensor
- ▶ Moves in straight lines or along obstacle boundaries




Bug Algorithm Assumptions

Environment

- ▶ Start location \vec{p}_{start} , goal location \vec{p}_{goal}
- ▶ Polygonal obstacles O_1, \dots, O_n

Robot

- ▶ Knows direction towards goal (absolute and relative)
- ▶ Knows the straight line distance to the goal
- ▶ Has no knowledge of obstacles (no map)
- ▶ Has a contact sensor
- ▶ Moves in straight lines or along obstacle boundaries
- ▶ Has memory to store distances and angles



media/video/demo_bug_direct_success.mp4


media/video/demo_bug_direct_fail.mp4

media/video/demo_bug_0_success.mp4

The Bug 0 algorithm

- 1: **while** not at goal **do**
- 2: move towards the goal
- 3: **if** hit an obstacle **then**
- 4: **while** not able to move towards the goal **do**
- 5: follow obstacle boundary CCW
- 6: **end while**
- 7: **end if**
- 8: **end while**

- ⊕ Algorithm solely depends on goal direction
- ⊖ The algorithm may fail...



Source: Lectures on Robotic Planning and Kinematics, Bullo and Smith, 2016


media/video/demo_bug_0_fail.mp4

The Bug 2 algorithm

```
1:  $\alpha \leftarrow$  goal direction
2: while not at goal do
3:   move towards the goal
4:   if hit an obstacle then
5:      $d_{\text{hit}} \leftarrow$  distance to goal
6:     while distance to goal  $\geq d_{\text{hit}}$  and
7:       goal direction  $\neq \alpha$  do
8:       follow obstacle boundary
9:     end while
10:  end if
11: end while
```

⊕ If \vec{p}_{goal} is reachable from \vec{p}_{start} , the algorithm will find a path

⊖ Not always better than Bug 1




Source: Lectures on Robotic Planning and Kinematics, Bullo and Smith, 2016

media/video/demo_bug_2.mp4


PART II

Spiking Neural Networks and the Neural Engineering Framework

Textbook Biological Model Neuron



Leaky Integrate and Fire Neuron

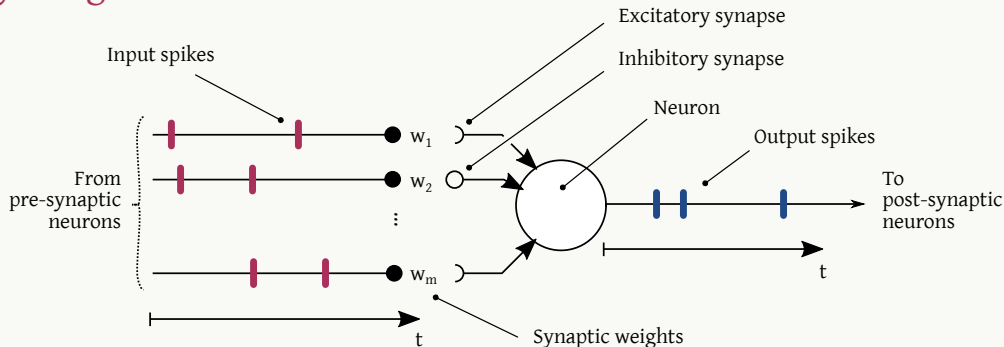


- ▶ Input spikes induce current J in cell body
- ▶ Cell body modelled as dynamical system in continuous time

$$\tau \dot{u}(t) = J - u(t)$$

$$u(t) \leftarrow 0 \text{ if } u(t) = 1$$

Leaky Integrate and Fire Neuron



- ▶ Input spikes induce current J in cell body
- ▶ Cell body modelled as dynamical system in continuous time



$$\tau \dot{u}(t) = J - u(t)$$

$$u(t) \leftarrow 0 \text{ if } u(t) = 1$$

⚠ Not your usual artificial neural network! Intrinsic dynamics! Spike noise!

The Neural Engineering Framework (NEF)

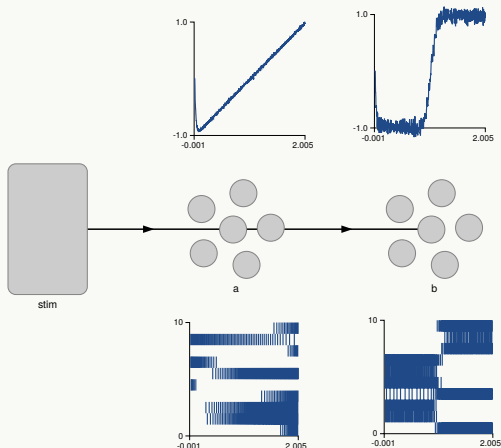
- ▶ Systematic way of building spiking neural networks
- ▶ *Principle 1:*
Populations of spiking neurons represent \vec{x}
- ▶ *Principle 2:*
Connections between populations compute $f(\vec{x}) = \vec{y}$
- ▶ *Principle 3:*
Self-connections implement dynamics $d/dt \vec{x} = g(\vec{x})$



Chris Eliasmith and
Charles H. Anderson,
Neural Engineering, 2003

Nengo

- Nengo: Python software implementation of the NEF, available on GitHub (Bekolay et al., Nengo: A Python tool for building large-scale functional brain models, Frontiers in Neuroinformatics, 2014)



```
1 import nengo
2 import numpy as np
3
4 model = nengo.Network()
5 with model:
6     stim = nengo.Node(
7         lambda t: t - 1)
8
9     a = nengo.Ensemble(
10         n_neurons=400, dimensions=1)
11     b = nengo.Ensemble(
12         n_neurons=400, dimensions=1)
13
14     nengo.Connection(stim, a)
15     nengo.Connection(a, b,
16         function=np.sign)
```


PART III

Implementation and Results

Simulator Environment


Body & Environment

- ▶ Discrete simulation
 $\Delta t = 10 \text{ ms}$
- ▶ Disk-shaped robot
- ▶ Robot slides along obstacles




Sensors

- ▶ Distance d , contact sensor
- ▶ Absolute and relative orientation (unit vectors $\vec{\alpha}_a, \vec{\alpha}_r$)
- ▶ Radar for obstacle boundary following



Motor System

- ▶ Non-holonomic 2-DOF drive
- ▶ Relative control vector $\vec{v} = (v_x, v_y)$



Basic Behaviours: Move Towards Goal & Follow Obstacle


Move Towards Goal

- ▶ Connect relative direction vector $\vec{\alpha}_r$ to the motor control output \vec{v}



Follow Obstacle

- ▶ Rotate radar vectors by $\approx 45^\circ$
- ▶ Weighted sum of radar vectors; smaller distance, more weight



Basic Behaviours: Move Towards Goal & Follow Obstacle


Move Towards Goal

- ▶ Connect relative direction vector $\vec{\alpha}_r$ to the motor control output \vec{v}



Follow Obstacle

- ▶ Rotate radar vectors by $\approx 45^\circ$
- ▶ Weighted sum of radar vectors; smaller distance, more weight



Basic Behaviours: Move Towards Goal & Follow Obstacle


Move Towards Goal

- ▶ Connect relative direction vector $\vec{\alpha}_r$ to the motor control output \vec{v}



Follow Obstacle

- ▶ Rotate radar vectors by $\approx 45^\circ$
- ▶ Weighted sum of radar vectors; smaller distance, more weight




media/video/demo_bug_follow.mp4

Bug 0: Reference Implementation

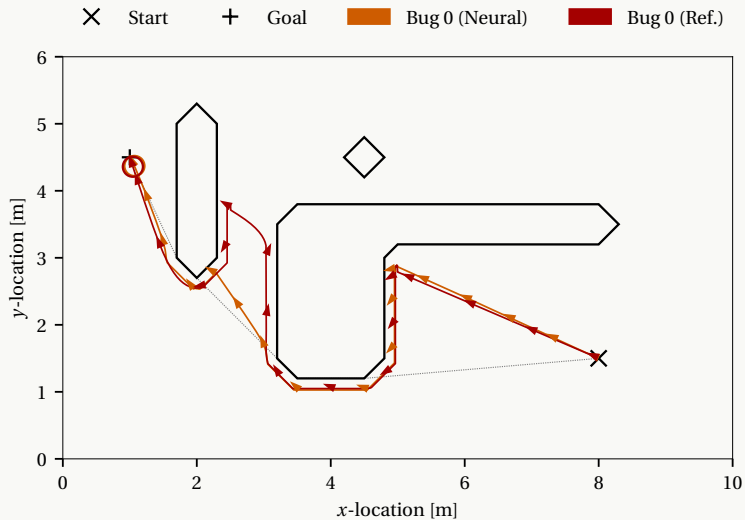
```
1  def behave(self, sensors, motor):
2      # State transition
3      if self.t_follow > 1.0:
4          self.follow_obstacle *= 0.5
5      elif sensors.hit_obstacle():
6          self.follow_obstacle += 1.0 * self.dt_coarse
7      else:
8          self.follow_obstacle -= self.follow_obstacle * self.dt_coarse
9
10     # Behaviour implementation
11     if self.follow_obstacle > 0.25:
12         common.follow_obstacle(sensors, motor, self.radius)
13         self.t_follow += self.dt_coarse
14     else:
15         self.t_follow = 0
16         common.move_towards_goal(sensors, motor)
```

Bug 0: Neural Network

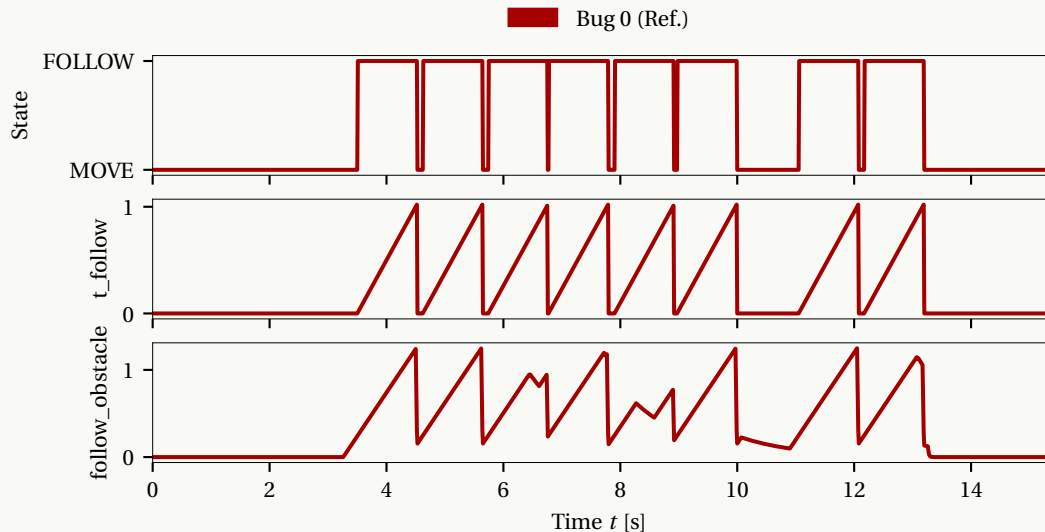


media/video/bug_0_comparison.mp4

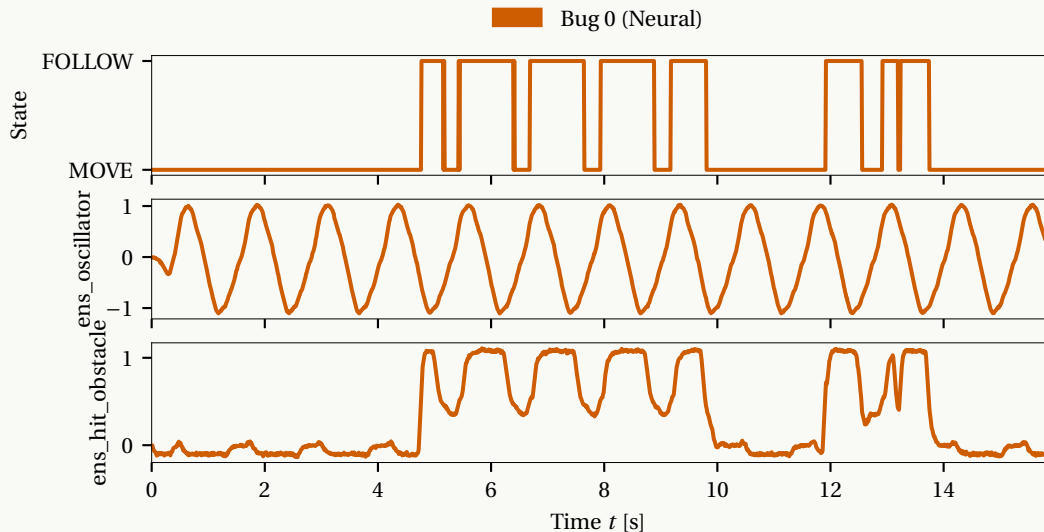
Bug 0: State over Time



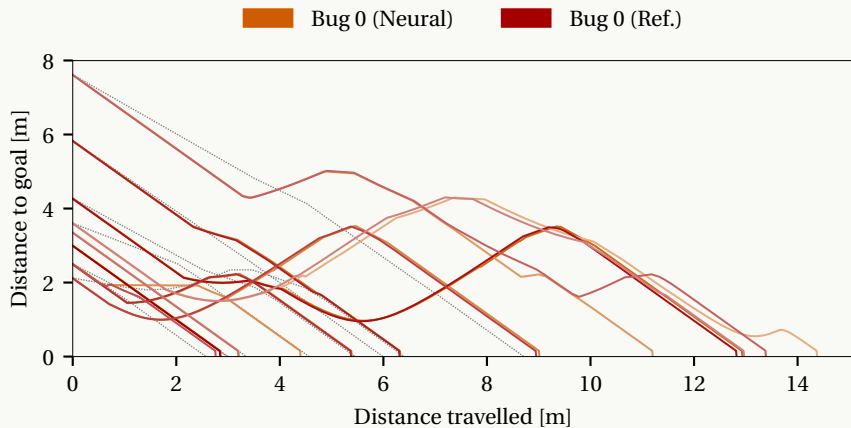
Bug 0: State over Time



Bug 0: State over Time



Bug 0: Distance Plot



Bug 2: Thinking About States

```
1:  $\alpha \leftarrow$  goal direction
2: while not at goal do
3:   move towards the goal
4:   if hit an obstacle then
5:      $d_{\text{hit}} \leftarrow$  distance to goal
6:     while distance to goal  $\geq d_{\text{hit}}$  and
7:       goal direction  $\neq \alpha$  do
8:       follow obstacle boundary
9:     end while
10:  end if
11: end while
```

► *Simulator/Neural network:*
Executes algorithm “tick-wise”

↪ Translate to state machine

Bug 2: Thinking About States

```
1:  $\alpha \leftarrow$  goal direction
2: while not at goal do
3:   move towards the goal
4:   if hit an obstacle then
5:      $d_{\text{hit}} \leftarrow$  distance to goal
6:     while distance to goal  $\geq d_{\text{hit}}$  and
7:       goal direction  $\neq \alpha$  do
8:       follow obstacle boundary
9:     end while
10:  end if
11: end while
```

► *Simulator/Neural network:*
Executes algorithm “tick-wise”

↪ Translate to state machine

States:

1. MEM_DIR: Memorize direction
2. MOVE: Move towards goal
3. MEM_DIST: Memorize distance
4. FOLLOW: Follow obstacle outline

Bug 2: Thinking About States

```
1:  $\alpha \leftarrow$  goal direction
2: while not at goal do
3:   move towards the goal
4:   if hit an obstacle then
5:      $d_{\text{hit}} \leftarrow$  distance to goal
6:     while distance to goal  $\geq d_{\text{hit}}$  and
7:       goal direction  $\neq \alpha$  do
8:       follow obstacle boundary
9:     end while
10:  end if
11: end while
```

► *Simulator/Neural network:*
Executes algorithm “tick-wise”

↪ Translate to state machine

States:

1. MEM_DIR: Memorize direction
2. MOVE: Move towards goal
3. MEM_DIST: Memorize distance
4. FOLLOW: Follow obstacle outline

State Transition Table:

$\epsilon \rightarrow 1, \quad 1 \rightarrow 2, \quad 2 \rightarrow 3, \quad 4 \rightarrow 2$

Bug 2: Thinking About States

```
1:  $\alpha \leftarrow$  goal direction
2: while not at goal do
3:   move towards the goal
4:   if hit an obstacle then
5:      $d_{\text{hit}} \leftarrow$  distance to goal
6:     while distance to goal  $\geq d_{\text{hit}}$  and
7:       goal direction  $\neq \alpha$  do
8:       follow obstacle boundary
9:     end while
10:  end if
11: end while
```

► *Simulator/Neural network:*
Executes algorithm “tick-wise”

↪ Translate to state machine

States:


1. MEM_DIR: Memorize direction
2. MOVE: Move towards goal
3. MEM_DIST: Memorize distance
4. FOLLOW: Follow obstacle outline

State Transition Table:

$\epsilon \rightarrow 1, \quad 1 \rightarrow 2, \quad 2 \rightarrow 3, \quad 4 \rightarrow 2$

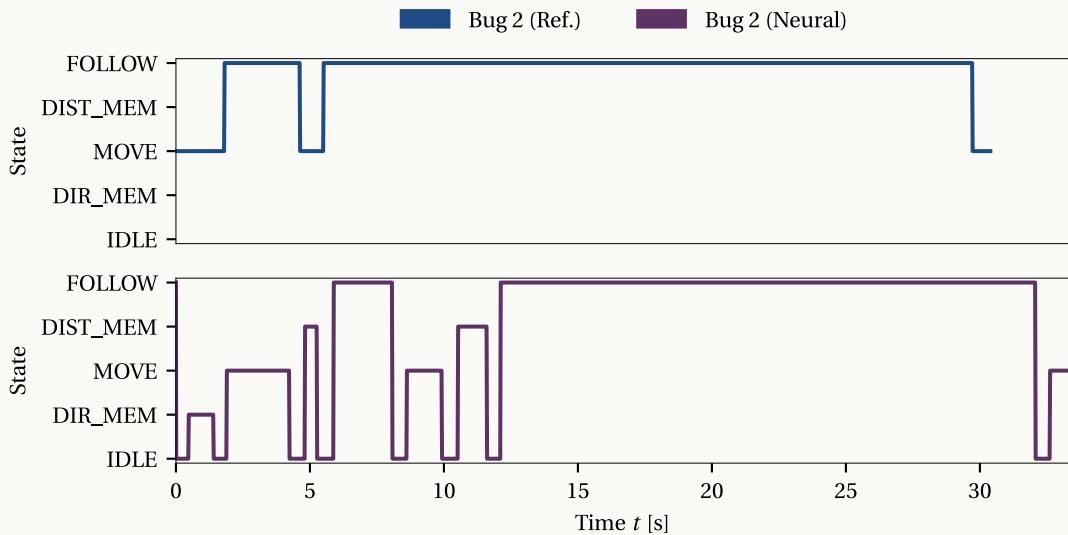
↪ How to implement this as a
spiking neural network?

Bug 2: State Transition Network




media/video/bug_2_comparison.mp4

Bug 2: State over Time



Bug 2: Distance Plot



Conclusion

Biological Plausibility

- ▶ Given the sensors, neurobiological implementation plausible, only few thousand neurons required
- ▶ Many of the sensors map to biological systems:
 - ▶ radar \leftrightarrow structure from motion
 - ▶ contact \leftrightarrow tactile sensors
 - ▶ distance and direction \leftrightarrow olfactory system in insects

⚠ Yet, the given implementation must not be seen as a biological model!

Conclusion

Biological Plausibility

- ▶ Given the sensors, neurobiological implementation plausible, only few thousand neurons required
- ▶ Many of the sensors map to biological systems:
 - ▶ radar \leftrightarrow structure from motion
 - ▶ contact \leftrightarrow tactile sensors
 - ▶ distance and direction \leftrightarrow olfactory system in insects

⚠ Yet, the given implementation must not be seen as a biological model!

Lessons

- ▶ Implementing the simulator: doing CG right is hard...
- ▶ Implementation of *Bug 0* as neural network quite trivial and works well.
- ▶ Implementation of *Bug 2* suffers from noise/chaotic behaviour.

Conclusion

Biological Plausibility

- ▶ Given the sensors, neurobiological implementation plausible, only few thousand neurons required
- ▶ Many of the sensors map to biological systems:
 - ▶ radar \leftrightarrow structure from motion
 - ▶ contact \leftrightarrow tactile sensors
 - ▶ distance and direction \leftrightarrow olfactory system in insects

⚠ Yet, the given implementation must not be seen as a biological model!

Lessons

- ▶ Implementing the simulator: doing CG right is hard...
- ▶ Implementation of *Bug 0* as neural network quite trivial and works well.
- ▶ Implementation of *Bug 2* suffers from noise/chaotic behaviour.

Challenges


- ▶ Plenty of magic constants in both reference and neural implementation
- ▶ State transition network

Thank you for your attention!

The Bug 1 algorithm

- 1: **while** not at goal **do**
- 2: move towards the goal
- 3: **if** hit an obstacle **then**
- 4: **at the same time**
- 5: circumnavigate the obstacle
- 6: track minimum distance to goal
- 7: follow boundary back to minimum
- 8: **end if**
- 9: **end while**

- ⊕ One can show that this algorithm is *complete*
- ⊖ Relatively complex, paths are far from optimal



Source: Lectures on Robotic Planning and Kinematics, Bullo and Smith, 2016

media/video/demo_bug_1.mp4

Spiking Neural Networks: Why Bother?

- ▶ Can implement linear algebra and dynamical systems in spiking neural networks

Spiking Neural Networks: Why Bother?


- ▶ Can implement linear algebra and dynamical systems in spiking neural networks
- ⇒ Why not just use the mathematical description?

Spiking Neural Networks: Why Bother?

- ▶ Can implement linear algebra and dynamical systems in spiking neural networks
- ⇒ Why not just use the mathematical description?

ENGINEERING

Use neuromorphic hardware




Spikey Neuromorphic Chip

Source: Electronic Visions Group, KIP Heidelberg

COMPUTATIONAL NEUROSCIENCE

Build biologically constrained models of cognition



Human Connectome. Data by Horn A. et al., 2014

Source: Wikimedia Commons